

X decides what you see. You don't.

WHO Founders & creators **building in public on X** — 10k–500k followers who can't just log off. X is their distribution, research, and network at once.

TODAY They doomscroll past ragebait, crypto shills, and ads-as-posts to reach the 3 posts that matter — or fight it with keyword mutes that don't understand anything.

WHY IT SUCKS The algorithm optimizes for engagement, not your time. The one place you can't quit is the one you least control.

Home



Rin Sato @rin_builds

Shipped the new onboarding flow today — cut signup drop-off by a third.

Hidden — ~~rage / outrage bait~~

X rage bait 81% [Show](#)

Hidden — ~~ads / promotions~~

X ad 94% [Show](#)



Jonas Meier @jmeier

Your best marketing channel is a changelog people actually want to read.

Hidden — ~~author @spamaccount~~

[Show](#)

An AI that re-ranks your feed to **your** rules — **on** **your device.**

A Chrome extension reads your timeline as it renders and hides everything you told it to, in **plain language** — "hide AI hype threads," "hide anything that's obviously an ad."

UNDERSTANDS, ≠ MATCHES

An LLM reads every post

Judges meaning, not keywords. Sarcasm, subtext, "ad dressed as a post" — caught.

ZERO SERVERS · ZERO COST

Runs 100% on-device

Chrome's built-in Gemini Nano. Nothing leaves the browser. \$0 marginal cost per user.

PLAIN LANGUAGE

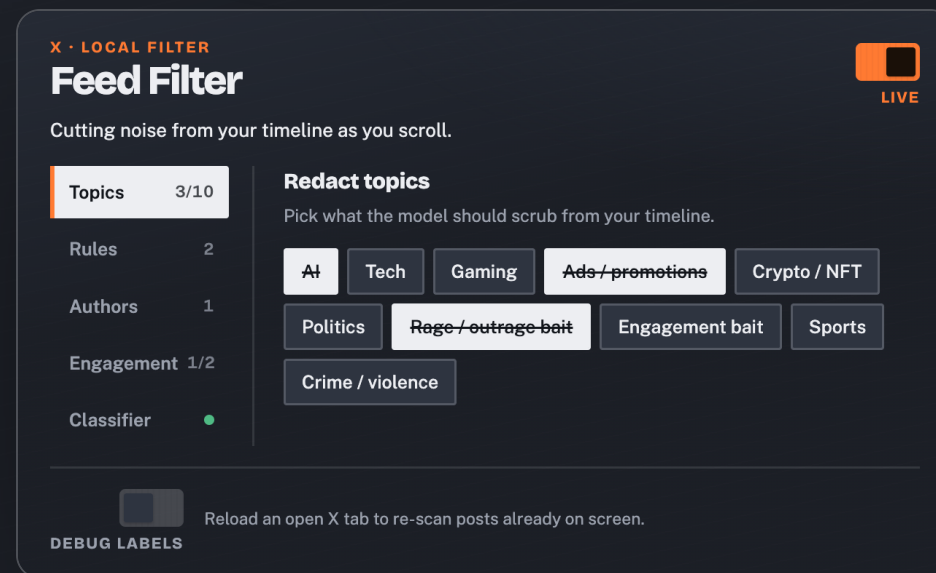
You write the rules

10 one-click presets + unlimited custom rules. No regex, no block-list babysitting.

YOU OWN RANKING

Not X, not us

Thread-aware, cached across scroll, fail-open. The algorithm just supplies raw posts.



The control panel — presets, rules, authors, on-device classifier

Home

 **Rin Sato** @rin_builds · 2h

Shipped the new onboarding flow today — cut signup drop-off by a third. Small copy changes made a bigger difference than the redesign did.

 14  6  128  3.9K ✓ kept 92% ER 3.8%

Hidden — ai ✗ ai 96% [Show](#)

 **Jonas Meier** @jmeier · 4h

Hot take: your best marketing channel is a changelog people actually want to read. Nobody subscribes to a roadmap.

 61  44  940  12.1K ✓ kept 88% HOT 8.4%

Hidden — rage / outrage bait ✗ rage / outrage bait 81% [Show](#)

Hidden — author @spamaccount [Show](#)

 **Anna Krüger** @akruger · 6h

Stuttgart in July: warm days, longer queues at the Markthalle, and every rooftop bar suddenly fully booked.

 9  2  71  1.6K ✓ kept 95% ER 4.9%

Same timeline. Your rules applied.

✓ kept 92% **Signal stays**
Real build-in-public posts pass, tagged with the model's confidence.

✗ ai 96% **Noise collapses in place**
Matched posts fold to a one-line receipt — reason + confidence, one click to reveal.

HOT 8.4% **Engagement rate, exposed**
(likes+replies+reposts) ÷ views on every post. A **Hot** badge flags real traction.

ER 3.8% **Nothing sent anywhere**
Every verdict on screen was computed locally, in the browser, in real time.

Would we use it ourselves? **Every day.**

We built it because **we** live on X and hate the noise. It's the tool we wanted — and it's already running on our own feeds.

Because inference is on-device, our cost to serve a user is **exactly \$0**. That means we can give the core away free, forever — no burn, no scaling wall.

\$0

marginal cost per user —
no servers to run

1 file

per new platform —
engine is already
platform-agnostic

Sell it to other creators? **That's the plan.**

FREE

On-device core

Presets, custom rules, blocklists. Free forever — creators post about tools they love. Built-in loop.

PRO

Rule sync + stronger models

Cross-device sync, saved rule packs, and a swap-in API model for heavier feeds.

TEAMS

Shared blocklists & presets

Agencies and creator teams standardize a clean feed across the whole roster.

EXPAND

LinkedIn · Reddit · more

Same engine, one new adapter each. X power-users are the beachhead, not the ceiling.

Why now

Chrome just shipped a real LLM inside the browser (Gemini Nano, v138). For the first time you can run a genuine classifier on every post — **free, private, no API bill.**

Two years ago this was a server product with a per-token cost that killed the unit economics. The on-device moment is **right now** — and trust in the algorithm has never been lower.

Why us

We shipped a **working on-device AI product** at a hackathon — not a mockup. Live classifier, thread-aware hiding, engagement analytics, a real settings UI.

We're the users: creators who read X all day and cared enough to rebuild its ranking from scratch in a weekend.

Full-time if it gets traction? Yes — the roadmap (sync, Pro, new platforms) is already scoped.

Take the ranking decision back from the algorithm — and never send a byte to do it.